

Bayesian ML

A practical approach

Verena, Filip

Motivation

Key issue of ML: overfitting

→ solution at hand: regularisation (dropout,...)

Became concern for us with Model-based Reinforcement Learning.

Disclaimer: only in the research phase. No successful application for RL yet.

Filip will show you another application.

Overfitting



Example regression:

NN is a probabilistic model $p(y | \mathbf{x}, \mathbf{w})$, p is a Gaussian distribution.

With training data $\mathcal{D} = \{x^{(i)}, y^{(i)}\}$ optimise \mathbf{w} by maximising likelihood function $p(\mathcal{D} | \mathbf{w}) = \prod_i p(y^{(i)} | \mathbf{x}^{(i)}, \mathbf{w}) \dots$ Maximum

Likelihood Estimate \rightarrow point estimate of \mathbf{w}

Better: regularisation

Would like to maximise $p(\mathbf{w} | \mathcal{D}) \dots$ maximum a posteriori estimate of \mathbf{w}


Do this with $p(\mathbf{w} | \mathcal{D}) \propto p(\mathcal{D} | \mathbf{w})p(\mathbf{w}) \dots$ prior distribution $p(\mathbf{w})$.

\rightarrow still only point estimate of \mathbf{w}

Example

Dataset is 32 points of noisy sinusoidal function.

```
18 def f(x, sigma):
19     epsilon = np.random.randn(*x.shape) * sigma
20     return 10 * np.sin(2 * np.pi * (x)) + epsilon
21
22 train_size = 32
23 noise = 1.0
24
25 X = np.linspace(-0.5, 0.5, train_size).reshape(-1, 1)
26 y = f(X, sigma=noise)
27 y_true = f(X, sigma=0.0)
28
```



Example

Simple feedforward regression model with regularisation with TensorFlow 2

```

1 import tensorflow.keras as tfk
2 from tensorflow.keras import initializers, activations
3 from tensorflow.keras.layers import Layer
4 import tensorflow_probability as tfp
5 import tensorflow as tf
6 tfd = tfp.distributions
7

```

```


41 batch_size = train_size
42 num_batches = train_size / batch_size
43
44
45
46 # Build model.
47 model = tf.keras.Sequential([
48     tfk.layers.Input(shape=(1,)),
49     tfk.layers.Dense(20, activation="relu"),
50     tfk.layers.Dropout(0.1),
51     tfk.layers.Dense(20, activation="relu"),
52     tfk.layers.Dropout(0.1),
53     tfk.layers.Dense(1)
54     ,tfk.layers.Dropout(0.1)
55 ])
56
57
58
59 from tensorflow.keras import callbacks, optimizers
60
61 def neg_log_likelihood(y_obs, y_pred, sigma=noise):
62     dist = tfp.distributions.Normal(loc=y_pred, scale=sigma)
63     return tfk.backend.sum(-dist.log_prob(y_obs))
64
65 model.compile(loss=neg_log_likelihood, optimizer=optimizers.Adam(lr=0.08), metrics=['mse'])
66 model.fit(X, y, batch_size=batch_size, epochs=1500, verbose=1);
67

```

Model

← Assume $p(y | \mathbf{x}, \mathbf{w})$ is Gaussian

Example




Test of model:

```
68
69 X_test = np.linspace(-1.5, 1.5, 1000).reshape(-1, 1)
70 y_pred_list = []
71
72 import tqdm
73
74 for i in tqdm.tqdm(range(500)):
75     y_pred = model.predict(X_test)
76     y_pred_list.append(y_pred)
77
78 y_preds = np.concatenate(y_pred_list, axis=1)
79
80 y_mean = np.mean(y_preds, axis=1)
81 y_sigma = np.std(y_preds, axis=1)
82
83 plt.plot(X_test, y_mean, 'r-', label='Predictive mean');
84 plt.scatter(X, y, marker='+', label='Training data')
85 plt.fill_between(X_test.ravel(),
86                  y_mean + 2 * y_sigma,
87                  y_mean - 2 * y_sigma,
88                  alpha=0.5, label='Epistemic uncertainty')
89 plt.title('Prediction')
90 plt.legend();
91
92 plt.show()
```

Example



Result:



Can we trust this? Need **epistemic uncertainty** on predictions.

Towards Variational Layers



With regularisation already assume distributions of \mathbf{w} ...how to extract uncertainty on predictions of NN?

→ Average predictions over ensemble of NNs weighted by posterior probabilities of their \mathbf{w}

$$p(y | \mathbf{x}, \mathcal{D}) = \int p(y | \mathbf{x}, \mathbf{w}) p(\mathbf{w} | \mathcal{D}) d\mathbf{w}$$

But we cannot write $p(\mathbf{w} | \mathcal{D})$ down.

→ Use variational distribution $q(\mathbf{w} | \theta)$ of known functional form.
Need to find θ by minimising KL divergence between $p(\mathbf{w} | \mathcal{D})$ and $q(\mathbf{w} | \theta)$

Towards Variational Layers




Resulting cost function of KL divergence minimisation can be arranged as:

$$\mathcal{F}(\mathcal{D}, \theta) = \underbrace{\mathbb{E}_{q(\mathbf{w}|\theta)} \log q(\mathbf{w} | \theta) - \mathbb{E}_{q(\mathbf{w}|\theta)} \log p(\mathbf{w})}_{\text{Data independent}} - \mathbb{E}_{q(\mathbf{w}|\theta)} \log p(\mathcal{D} | \mathbf{w})$$

Data independent

See ELBO in <https://medium.com/tensorflow/regression-with-probabilistic-layers-in-tensorflow-probability-e46ff5d37baf>

Constructing a Variational Layer yourself



Assume Gaussian as variational posterior for \mathbf{w}

$$\rightarrow \theta = (\mu, \sigma)$$


Assume non-trainable prior:

$$p(\mathbf{w}) = \pi \mathcal{N}(\mathbf{w} \mid 0, \sigma_1) + (1 - \pi) \mathcal{N}(\mathbf{w} \mid 0, \sigma_2)$$

Basic steps:

- define your trainable parameters in `build()`
- define in `call()` how to evaluate your layer given `input`

Constructing a Variational Layer yourself



Example:

```
12
13 class MyDenseVariational(Layer):
14     def __init__(self,
15                 units,
16                 kl_weight,
17                 activation=None,
18                 prior_sigma_1=1.5,
19                 prior_sigma_2=0.1,
20                 prior_pi=0.5, **kwargs):
21         self.units = units
22         self.kl_weight = kl_weight
23         self.activation = activations.get(activation)
24         self.prior_sigma_1 = prior_sigma_1
25         self.prior_sigma_2 = prior_sigma_2
26         self.prior_pi_1 = prior_pi
27         self.prior_pi_2 = 1.0 - prior_pi
28         self.init_sigma = np.sqrt(self.prior_pi_1 * self.prior_sigma_1 ** 2 +
29                                   self.prior_pi_2 * self.prior_sigma_2 ** 2)
30
31         super().__init__(**kwargs)
32
33     def compute_output_shape(self, input_shape):...
34
35     def build(self, input_shape):...
36
37     def call(self, inputs, **kwargs):...
38
39     def kl_loss(self, w, mu, sigma):...
40
41     def log_prior_prob(self, w):...
```


Constructing a Variational Layer yourself




For the details of `build()` and `call()`:

```
36 def build(self, input_shape):
37     self.kernel_mu = self.add_weight(name='kernel_mu',
38                                     shape=(input_shape[1], self.units),
39                                     initializer=initializers.RandomNormal(stddev=self.init_sigma),
40                                     trainable=True)
41     self.bias_mu = self.add_weight(name='bias_mu',
42                                   shape=(self.units,),
43                                   initializer=initializers.RandomNormal(stddev=self.init_sigma),
44                                   trainable=True)
45     self.kernel_rho = self.add_weight(name='kernel_rho',
46                                     shape=(input_shape[1], self.units),
47                                     initializer=initializers.constant(0.0),
48                                     trainable=True)
49     self.bias_rho = self.add_weight(name='bias_rho',
50                                   shape=(self.units,),
51                                   initializer=initializers.constant(0.0),
52                                   trainable=True)
53     super().build(input_shape)
54
55 def call(self, inputs, **kwargs):
56     kernel_sigma = tf.math.softplus(self.kernel_rho)
57     kernel = self.kernel_mu + kernel_sigma * tf.random.normal(self.kernel_mu.shape)
58
59     bias_sigma = tf.math.softplus(self.bias_rho)
60     bias = self.bias_mu + bias_sigma * tf.random.normal(self.bias_mu.shape)
61
62     self.add_loss(self.kl_loss(kernel, self.kernel_mu, kernel_sigma) +
63                 self.kl_loss(bias, self.bias_mu, bias_sigma))
64
65     return self.activation(tfk.backend.dot(inputs, kernel) + bias)
66
67 def kl_loss(self, w, mu, sigma):
68     variational_dist = tfp.distributions.Normal(mu, sigma)
69     return self.kl_weight * tfk.backend.sum(variational_dist.log_prob(w) - self.log_prior_prob(w))
70
```


Using MyDenseVariational



```
105
106 kl_weight = 1.0 / num_batches
107 prior_params = {
108     'prior_sigma_1': 1.5,
109     'prior_sigma_2': 0.1,
110     'prior_pi': 0.5
111 }
112
113 x_in = Input(shape=(1,))
114 x = MyDenseVariational(20, kl_weight, **prior_params, activation='relu')(x_in)
115 x = MyDenseVariational(20, kl_weight, **prior_params, activation='relu')(x)
116 x = MyDenseVariational(1, kl_weight, **prior_params)(x)
117
118 model = Model(x_in, x)
119
```



Using TF2 DenseVariational Layer



```
69
70 # Build model.
71 model = tf.keras.Sequential([
72     tf.keras.layers.Input(shape=(1,)),
73     tfp.layers.DenseVariational(units=20,
74                                 make_posterior_fn=posterior_mean_field,
75                                 make_prior_fn=prior_trainable,
76                                 kl_weight=kl_weight,
77                                 activation='relu'),
78     tfp.layers.DenseVariational(units=20,
79                                 make_posterior_fn=posterior_mean_field,
80                                 make_prior_fn=prior_trainable,
81                                 kl_weight=kl_weight,
82                                 activation='relu'),
83     tfp.layers.DenseVariational(units=1,
84                                 make_posterior_fn=posterior_mean_field,
85                                 make_prior_fn=prior_trainable,
86                                 kl_weight=kl_weight)
87 ])
88
```

```
# Specify the surrogate posterior over `keras.layers.Dense` `kernel` and `bias`.
def posterior_mean_field(kernel_size, bias_size=0, dtype=None):
    n = kernel_size + bias_size
    c = np.log(np.exp(1.))
    return tf.keras.Sequential([
        tfp.layers.VariableLayer(2 * n, dtype=dtype),
        tfp.layers.DistributionLambda(lambda t: tfd.Independent(
            tfd.Normal(loc=t[..., :n],
                       scale=1e-5 + tf.nn.softplus(c+t[..., n:])),
            reinterpreted_batch_ndims=1)),
    ])

```


```
# Specify the prior over `keras.layers.Dense` `kernel` and `bias`.
def prior_trainable(kernel_size, bias_size=0, dtype=None):
    n = kernel_size + bias_size
    return tf.keras.Sequential([
        tfp.layers.VariableLayer(n, dtype=dtype),
        tfp.layers.DistributionLambda(lambda t: tfd.Independent(
            tfd.Normal(loc=t, scale=1),
            reinterpreted_batch_ndims=1)),
    ])

```

Using TF2 DenseVariational Layer



Result:



DenseFlipout

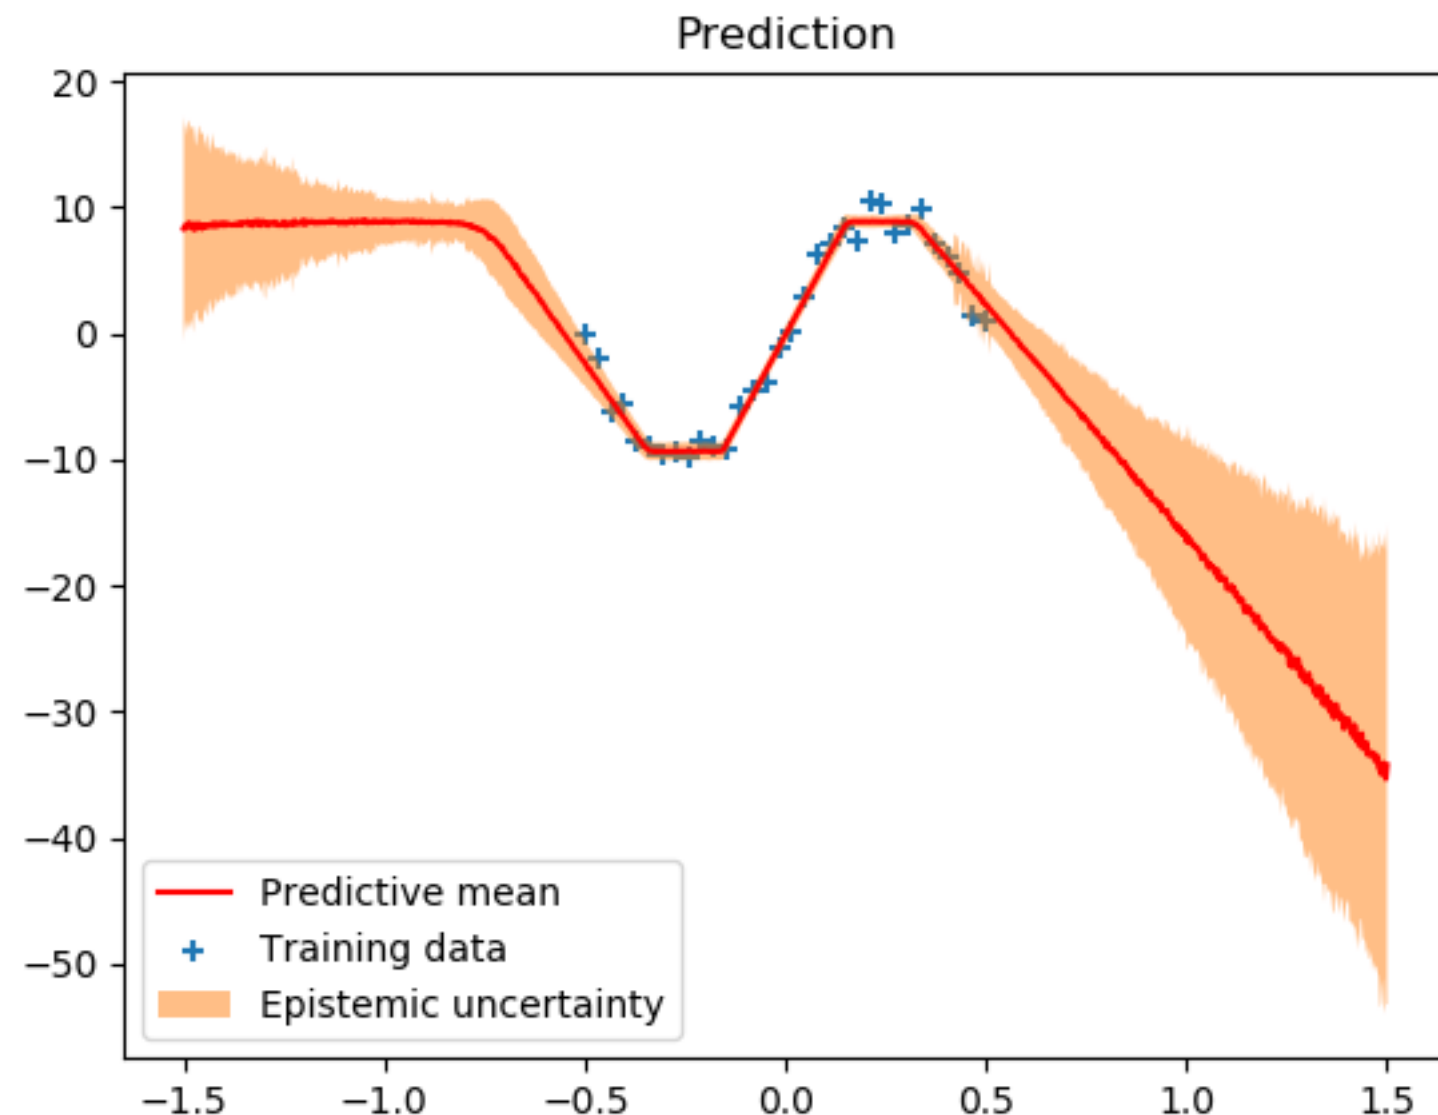


Assume distribution of weights and biases.

→ prior and posterior distributions for those to be defined.

Use MonteCarlo approach for sampling weights and biases and integrate over it.

```
# Build model.  
model = tf.keras.Sequential([  
    tf.keras.layers.Input(shape=(1,)),  
    tfp.layers.DenseFlipout(20, activation="relu"),  
    tfp.layers.DenseFlipout(20, activation="relu"),  
    tfp.layers.DenseFlipout(1)  
)
```



Further reading...



<https://brendanhasz.github.io/2019/07/23/bayesian-density-net.html>

<https://medium.com/tensorflow/regression-with-probabilistic-layers-in-tensorflow-probability-e46ff5d37baf>

<https://github.com/krasserm/bayesian-machine-learning>